

PR_HW3

January 26, 2020

1 Pattern Recognition - HW3 - Mohammad Doosti Lakhani - 98722278

Index:

1. Import Libraries
2. Loading Data
3. Pretrained Feature Extraction
4. Feature Reduction
 1. PCA
 1. 2D PCA
 2. Visualization
 3. 3D PCA
 4. Visualization
 2. TSN-E
 1. 2D TSN-E
 2. Visualization
 3. 3D TSN-E
 4. Visualization
5. Train Test Split
6. Train Models
 1. PCA
 1. SVC Linear - PCA 2D
 2. SVC Linear - PCA 3D
 3. SVC RBF - PCA 2D
 4. SVC RBF - PCA 3D
 2. TSN-E
 1. SVC Linear - TSN-E 2D
 2. SVC Linear - TSN-E 3D
 3. SVC RBF - TSN-E 2D
 4. SVC RBF - TSN-E 3D
 3. Fully Connected

1.1 1 Import Libraries

```
In [1]: %tensorflow_version 1.x

from __future__ import print_function

import pickle

import seaborn as sns
import matplotlib.pyplot as plt
import PIL
from mpl_toolkits.mplot3d import Axes3D

import pandas as pd
import numpy as np

from sklearn.svm import SVC, LinearSVC
from sklearn.svm import base

from sklearn.metrics import confusion_matrix
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE

import keras
from keras import backend as K
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers import Adam
from keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint, LearningRateScheduler
from keras.layers import Dense, Conv2D, BatchNormalization, Activation
from keras.layers import AveragePooling2D, Input, Flatten
from keras.regularizers import l2
from keras.models import Model
from keras.utils import to_categorical
from keras.applications.vgg19 import VGG19

/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:144: FutureWarning: The sklearn.svm.SVC
  warnings.warn(message, FutureWarning)
Using TensorFlow backend.
```

1.2 2 Loading Data

You need to put your data in the following structure:

```
data/train/
    /class1
    /class2
    /...
data/test/
```

```
/class1  
/class2  
/...
```

```
In [ ]: from google.colab import drive  
drive.mount('/content/drive')  
!unrar x drive/My\ Drive/IUST-PR/HW3/data.rar /content
```

```
In [17]: # loading data  
batch_size = 64  
n_classes = 3  
img_size = (224, 224)  
  
datagen = ImageDataGenerator()  
train_iterator = datagen.flow_from_directory('data/train/', class_mode='categorical',  
                                             batch_size=batch_size, target_size=img_s  
                                             batch_size=batch_size, target_size=img_s  
# only for NN model  
test_iterator = datagen.flow_from_directory('data/test/', class_mode='categorical',  
                                            batch_size=batch_size, target_size=img_s  
n = 541  
x = np.empty((n, img_size[0], img_size[1], n_classes))  
y = np.empty((n, n_classes))  
for idx in range(len(train_iterator)):  
    batchx = train_iterator[idx][0]  
    batchy = train_iterator[idx][1]  
    if batchx.shape[0] < batch_size:  
        x[idx*batch_size:] = batchx  
        y[idx*batch_size:] = batchy  
    else:  
        x[idx*batch_size:(idx+1)*batch_size] = batchx  
        y[idx*batch_size:(idx+1)*batch_size] = batchy  
  
n_test = 31  
y_test = np.empty((n_test, n_classes))  
for idx in range(len(test_iterator)):  
    batchy = test_iterator[idx][1]  
    if batchx.shape[0] < batch_size:  
        y_test[idx*batch_size:] = batchy  
    else:  
        y_test[idx*batch_size:(idx+1)*batch_size] = batchy  
  
print(x.shape)  
print(y.shape)  
print(y_test.shape)
```

Found 541 images belonging to 3 classes.

```
Found 31 images belonging to 3 classes.  
(541, 224, 224, 3)  
(541, 3)  
(31, 3)
```

1.3 3 Pretrained Feature Extraction

When a CNN model is trained on a large dataset that can represent images similar to our custom dataset, has features that are common in first layers and more specific in the last layers. So we remove fully connected layers and use convolutional layers as the features extracted from ImageNet dataset where it is the superset of our dataset. These features even though are in low dimensional manner w.r.t. input images, they contain much more usefull information about images.

```
In [18]: # feature extraction
```

```
model = VGG19(include_top=False, weights='imagenet', input_tensor=None, input_shape=(  
features = model.predict_generator(train_iterator, steps=9)  
features_test = model.predict_generator(test_iterator, steps=1)  
print('NN data train shape:', features.shape)  
print('NN data test shape:', features_test.shape)
```

```
NN data train shape: (541, 7, 7, 512)
```

```
NN data test shape: (31, 7, 7, 512)
```

1.4 4 Feature Reduction

In this step, we have 77512 features to feed to our SVM models to train. But we want to reduce this info and use only 2 and 3 feature that represent our whole features.

So we use PCA and TSN-E to reduce to 2 and 3 features.

1. PCA
 1. 2D PCA
 2. Visualization
 3. 3D PCA
 4. Visualization
2. TSN-E
 1. 2D TSN-E
 2. Visualization
 3. 3D TSN-E
 4. Visualization

1.4.1 4.A PCA

1. 2D PCA
2. Visualization
3. 3D PCA
4. Visualization

4.A.a 2D PCA

```
In [33]: # 2D PCA
pca_2d = PCA(n_components=2)
pca_2d.fit(features.reshape(len(features), -1))
pca_2d_features = pca_2d.transform(features.reshape(len(features), -1))
print(pca_2d_features.shape)
```

(541, 2)

4.A.b Visualization


```
In [6]: colors = np.argmax(y, axis=1)
fig, ax = plt.subplots()
x_ = pca_2d_features[:,0]
y_ = pca_2d_features[:,1]
legend_ = ['blue', 'green', 'red']
for idx, color in enumerate(['car', 'bike', 'person']):
    area = (15 * np.random.rand(len(x)))*2
    print(x_[colors==idx].shape, y_[colors==idx].shape, colors[colors==idx].shape)
    ax.scatter(x_[colors==idx], y_[colors==idx], s=area, c=legend_[idx], alpha=0.5, label=idx)

ax.legend()
plt.show()
```

(165,) (165,) (165,)

(169,) (169,) (169,)

(207,) (207,) (207,)



4.A.c 3D PCA

```
In [34]: # 3D PCA
pca_3d = PCA(n_components=3)
pca_3d.fit(features.reshape(len(features), -1))
pca_3d_features = pca_3d.transform(features.reshape(len(features), -1))
print(pca_3d_features.shape)
```


(541, 3)

4.A.d Visualization

```
In [35]: colors = np.argmax(y, axis=1)
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
x_ = pca_3d_features[:, 0]
y_ = pca_3d_features[:, 1]
z_ = pca_3d_features[:, 2]
legend_ = ['blue', 'green', 'red']
for idx, color in enumerate(['car', 'bike', 'person']):
    area = (9 * np.random.rand(len(x)))**2
    print(x_[colors==idx].shape, y_[colors==idx].shape, colors[colors==idx].shape)
    ax.scatter(x_[colors==idx], y_[colors==idx], z_[colors==idx], s=area, c=legend_[i

ax.legend()
plt.show()
```

(165,) (165,) (165,)
(169,) (169,) (169,)
(207,) (207,) (207,)



1.4.2 4.B TSN-E

1. 2D TSN-E
2. Visualization
3. 3D TSN-E
4. Visualization

4.B.a 2D TSN-E

```
In [38]: # TSNE 2D
tsne_2d = TSNE(n_components=2, perplexity=30, learning_rate=200)
tsne_2d_features = tsne_2d.fit_transform(features.reshape(len(features), -1))
print(tsne_2d_features.shape)
```

(541, 2)

4.B.b Visualization

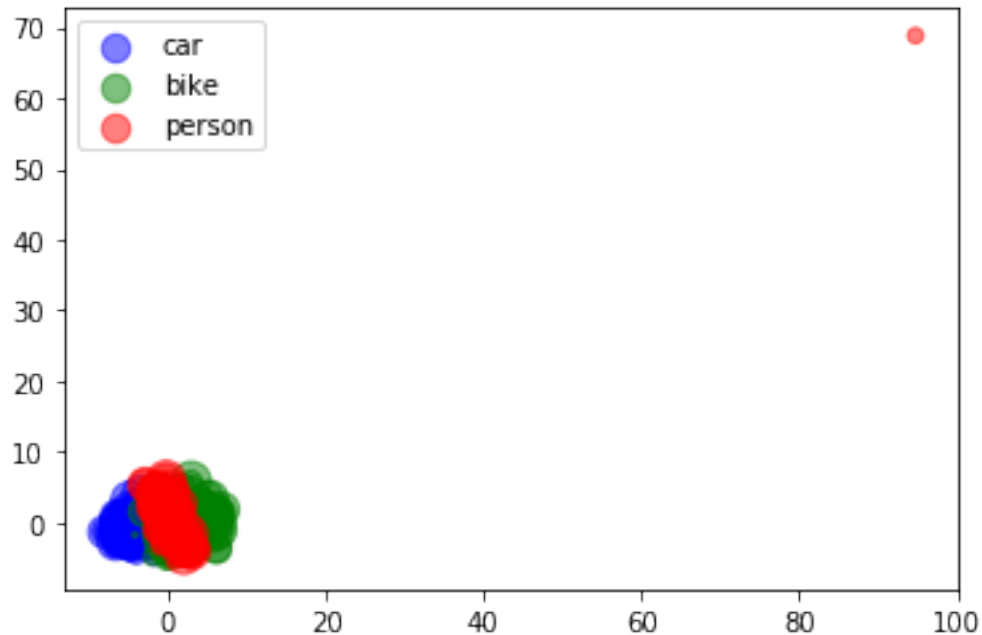
```

In [40]: colors = np.argmax(y, axis=1)
fig, ax = plt.subplots()
x_ = tsne_2d_features[:,0]
y_ = tsne_2d_features[:,1]
legend_ = ['blue', 'green', 'red']
for idx, color in enumerate(['car', 'bike', 'person']):
    area = (15 * np.random.rand(len(x)))*2
    print(x_[colors==idx].shape, y_[colors==idx].shape, colors[colors==idx].shape)
    ax.scatter(x_[colors==idx], y_[colors==idx], s=area, c=legend_[idx], alpha=0.5, label=idx)

ax.legend()
plt.show()

(165,) (165,) (165,)
(169,) (169,) (169,)
(207,) (207,) (207,)

```



4.B.c 3D TSN-E

```

In [37]: # TSNE 3D
tsne_3d = TSNE(n_components=3, perplexity=30, learning_rate=200)
tsne_3d_features = tsne_3d.fit_transform(features.reshape(len(features), -1))
print(tsne_3d_features.shape)

(541, 3)


```


4.B.d Visualization

```
In [39]: colors = np.argmax(y, axis=1)
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
x_ = tsne_3d_features[:, 0]
y_ = tsne_3d_features[:, 1]
z_ = tsne_3d_features[:, 2]
legend_ = ['blue', 'green', 'red']
for idx, color in enumerate(['car', 'bike', 'person']):
    area = (9 * np.random.rand(len(x)))**2
    print(x_[colors==idx].shape, y_[colors==idx].shape, colors[colors==idx].shape)
    ax.scatter(x_[colors==idx], y_[colors==idx], z_[colors==idx], s=area, c=legend_[i])

ax.legend()
plt.show()

(165,) (165,) (165,)
(169,) (169,) (169,)
(207,) (207,) (207,)
```



1.5 5 Train Test Split

For validating our trained model we need to separate a small proportion of our data to propose to our trained model where those have not been seen at all.

```
In [41]: # train test split
         from sklearn.model_selection import train_test_split

         pca_2d_x_train, pca_2d_x_test, pca_2d_y_train, pca_2d_y_test = train_test_split(pca_2d_x,
                                                                                       np.array(pca_2d_y),
                                                                                       test_size=0.1,
                                                                                       random_state=42)

         pca_3d_x_train, pca_3d_x_test, pca_3d_y_train, pca_3d_y_test = train_test_split(pca_3d_x,
                                                                                       np.array(pca_3d_y),
                                                                                       test_size=0.1,
                                                                                       random_state=42)

         tsne_2d_x_train, tsne_2d_x_test, tsne_2d_y_train, tsne_2d_y_test = train_test_split(tsne_2d_x,
                                                                                             np.array(tsne_2d_y),
                                                                                             test_size=0.1,
                                                                                             random_state=42)

         tsne_3d_x_train, tsne_3d_x_test, tsne_3d_y_train, tsne_3d_y_test = train_test_split(tsne_3d_x,
                                                                                             np.array(tsne_3d_y),
                                                                                             test_size=0.1,
                                                                                             random_state=42)

         print('PCA 2D features train size:',pca_2d_x_train.shape, '-----','PCA 2D features test size:',pca_2d_x_test.shape)
         print('PCA 3D features train size:',pca_3d_x_train.shape, '-----','PCA 3D features test size:',pca_3d_x_test.shape)
         print('TSNE 2D features train size:',tsne_2d_x_train.shape, '-----','TSNE 2D features test size:',tsne_2d_x_test.shape)
         print('TSNE 3D features train size:',tsne_3d_x_train.shape, '-----','TSNE 3D features test size:',tsne_3d_x_test.shape)
```

```
PCA 2D features train size: (486, 2) ----- PCA 2D features test size (55, 2)
PCA 3D features train size: (486, 3) ----- PCA 3D features test size (55, 3)
TSNE 2D features train size: (486, 2) ----- TSNE 2D features test size (55, 2)
TSNE 3D features train size: (486, 3) ----- TSNE 3D features test size (55, 3)
```

1.6 6 Train Models

1. PCA 1. SVC Linear - PCA 2D 2. SVC Linear - PCA 3D 3. SVC RBF - PCA 2D 4. SVC RBF - PCA 3D
2. TSN-E
 1. SVC Linear - TSN-E 2D
 2. SVC Linear - TSN-E 3D
 3. SVC RBF - TSN-E 2D
 4. SVC RBF - TSN-E 3D
3. Fully Connected

```
In [ ]: # SVM

class SVM_MODELS:
    """
    Enum of possible models
```

```

        """
        SVC = 1
        LINEAR_SVC = 2

# %% functions
def build_svm(model_type: int, x: np.ndarray, y: np.ndarray, svc_kernel: str = None, save: bool = False,
              path: str = None, **kwargs) -> base:
    """
    Trains a SVM model

    :param model_type: The kernel of SVM model (see `SVM_MODELS` class)
    :param svc_kernel: The possible kernels for `SVC` model (It must be one of linear, rbf, poly)
    :param x: Features in form of numpy ndarray
    :param y: Labels in form of numpy ndarray
    :param save: Whether save trained model on disc or not
    :param path: Path to save fitted model
    :param kwargs: A dictionary of other optional arguments of models in format of {'key': value}
    :return: A trained SVM model
    """
    if model_type == SVM_MODELS.SVC:
        model = SVC(kernel=svc_kernel, **kwargs)
    elif model_type == SVM_MODELS.LINEAR_SVC:
        model = LinearSVC(**kwargs)
    else:
        raise Exception('Model {} is not valid'.format(model_type))

    model.fit(x, y)

    if save:
        if path is None:
            path = ''
        pickle.dump(model, open(path + model.__module__ + '.model', 'wb'))
    return model

```

1.6.1 6.A PCA

1. SVC Linear - PCA 2D
2. SVC Linear - PCA 3D
3. SVC RBF - PCA 2D
4. SVC RBF - PCA 3D

6.A.a SVC Linear - PCA 2D

```

In [43]: # SVC LINEAR PCA 2D
pca_2d_svc_linear = build_svm(model_type=SVM_MODELS.LINEAR_SVC, x=pca_2d_x_train, y=pca_2d_y_train,
                               max_iter=200000, path='models/pca_2d_', multi_class='ovr')

```


```

print('PCA 2D - SVC LINEAR - Test accuracy', pca_2d_svc_linear.score(pca_2d_x_test, p
cm = confusion_matrix(pca_2d_svc_linear.predict(pca_2d_x_test), pca_2d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC LINEAR - PCA')
plt.show()

```

/usr/local/lib/python3.6/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear
"the number of iterations.", ConvergenceWarning)

PCA 2D - SVC LINEAR - Test accuracy 0.9090909090909091



6.A.b SVC Linear - PCA 3D

```


In [44]: # SVC LINEAR PCA 3D
pca_3d_svc_linear = build_svm(model_type=SVM_MODELS.LINEAR_SVC, x=pca_3d_x_train, y=p
max_iter=200000, path='models/pca_3d_', multi_class='ovr')
print('PCA 3D - SVC LINEAR - Test accuracy', pca_3d_svc_linear.score(pca_3d_x_test, p
cm = confusion_matrix(pca_3d_svc_linear.predict(pca_3d_x_test), pca_3d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)

```

```
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC LINEAR - PCA')
plt.show()
```

/usr/local/lib/python3.6/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear
"the number of iterations.", ConvergenceWarning)

PCA 3D - SVC LINEAR - Test accuracy 0.9272727272727272




6.A.c SVC RBF - PCA 2D

In [45]: # SVC RBF PCA 2D

```
pca_2d_svc_rbf = build_svm(model_type=SVM_MODELS.SVC, x=pca_2d_x_train, y=pca_2d_y_train,
                             path='models/pca_2d_', svc_kernel='rbf')

print('PCA 2D - SVC RBF - Test accuracy', pca_2d_svc_rbf.score(pca_2d_x_test, pca_2d_y_test))
cm = confusion_matrix(pca_2d_svc_rbf.predict(pca_2d_x_test), pca_2d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC RBF - PCA')
plt.show()
```

PCA 2D - SVC RBF - Test accuracy 0.9090909090909091




6.A.d SVC RBF - PCA 3D

In [46]: # SVC RBF PCA 3D

```
pca_3d_svc_rbf = build_svm(model_type=SVM_MODELS.SVC, x=pca_3d_x_train, y=pca_3d_y_train,
                             path='models/pca_3d_', svc_kernel='rbf')
```

```
print('PCA 3D - SVC RBF - Test accuracy', pca_3d_svc_rbf.score(pca_3d_x_test, pca_3d_y_test))
cm = confusion_matrix(pca_3d_svc_rbf.predict(pca_3d_x_test), pca_3d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC RBF - PCA')
plt.show()
```

PCA 3D - SVC RBF - Test accuracy 0.9272727272727272



1.6.2 6.B TSN-E


1. SVC Linear - TSN-E 2D
2. SVC Linear - TSN-E 3D
3. SVC RBF - TSN-E 2D
4. SVC RBF - TSN-E 3D

6.B.a SVC Linear - TSN-E 2D

In [47]: # SVC LINEAR TSNE 2D

```
tsne_2d_svc_linear = build_svm(model_type=SVM_MODELS.LINEAR_SVC, x=tsne_2d_x_train, y=
                                max_iter=200000, path='models/tsne_2d_', multi_class='ovr')
print('tsne 2D - SVC LINEAR - Test accuracy', tsne_2d_svc_linear.score(tsne_2d_x_test,
                                tsne_2d_y_test))
cm = confusion_matrix(tsne_2d_svc_linear.predict(tsne_2d_x_test), tsne_2d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC LINEAR - tsne')
plt.show()
```

tsne 2D - SVC LINEAR - Test accuracy 0.7636363636363637




6.B.b SVC Linear - TSN-E 3D

In [48]: # SVC LINEAR TSNE 3D

```
tsne_3d_svc_linear = build_svm(model_type=SVM_MODELS.LINEAR_SVC, x=tsne_3d_x_train, y=
                                max_iter=200000, path='models/tsne_3d_', multi_class='ovr')
print('tsne 3D - SVC LINEAR - Test accuracy', tsne_3d_svc_linear.score(tsne_3d_x_test,
                                tsne_3d_y_test))
cm = confusion_matrix(tsne_3d_svc_linear.predict(tsne_3d_x_test), tsne_3d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC LINEAR - tsne')
plt.show()
```

/usr/local/lib/python3.6/dist-packages/sklearn/svm/_base.py:947: ConvergenceWarning: Liblinear
"the number of iterations.", ConvergenceWarning)

tsne 3D - SVC LINEAR - Test accuracy 0.8727272727272727




6.B.c SVC RBF - TSN-E 2D

In [49]: # SVC RBF tsne 2D

```
tsne_2d_svc_rbf = build_svm(model_type=SVM_MODELS.SVC, x=tsne_2d_x_train, y=tsne_2d_y_train,
                             path='models/tsne_2d_', svc_kernel='rbf')

print('tsne 2D - SVC RBF - Test accuracy', tsne_2d_svc_rbf.score(tsne_2d_x_test, tsne_2d_y_test))
cm = confusion_matrix(tsne_2d_svc_rbf.predict(tsne_2d_x_test), tsne_2d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC RBF - tsne')
plt.show()
```

tsne 2D - SVC RBF - Test accuracy 0.9272727272727272




6.B.d SVC RBF - TSN-E 3D

In [50]: # SVC RBF tsne 3D

```
tsne_3d_svc_rbf = build_svm(model_type=SVM_MODELS.SVC, x=tsne_3d_x_train, y=tsne_3d_y_train,
                             path='models/tsne_3d_', svc_kernel='rbf')

print('tsne 3D - SVC RBF - Test accuracy', tsne_3d_svc_rbf.score(tsne_3d_x_test, tsne_3d_y_test))
cm = confusion_matrix(tsne_3d_svc_rbf.predict(tsne_3d_x_test), tsne_3d_y_test)
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('SVC RBF - tsne')
plt.show()
```

tsne 3D - SVC RBF - Test accuracy 0.8909090909090909



1.6.3 6.C Fully Connected

```
In [27]: def fc(input_shape, n_classes, print_summary=True):
    inputs = Input(shape=input_shape)
    x = Flatten(name='flatten')(inputs)
    x = Dense(64, activation='relu', name='fc1')(x)
    x = Dense(64, activation='relu', name='fc2')(x)
    x = Dense(n_classes, activation='softmax', name='predictions')(x)
    model = Model(inputs=inputs, outputs=x)
    if print_summary:
        model.summary()
    return model

fc_model = fc(features.shape[1:], n_classes)
fc_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

import os
def prepare_directory(model_type='vgg19'):
    save_dir = os.path.join(os.getcwd(), 'models')
    model_name = '%s_model.{epoch:03d}.h5' % model_type
    if not os.path.isdir(save_dir):
        os.makedirs(save_dir)
    filepath = os.path.join(save_dir, model_name)
    return filepath
```

```

filepath = prepare_directory()
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_acc', verbose=0, save_best_only=True,
                             callbacks = [checkpoint])

fc_model.fit(features, y,
              validation_split=0.1, epochs=50,
              verbose=1, workers=1, callbacks=callbacks)
score = fc_model.evaluate(features, y, batch_size=64)
print(fc_model.metrics_names)
print(score)

```

Model: "model_5"

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 64)	1605696
fc2 (Dense)	(None, 64)	4160
predictions (Dense)	(None, 3)	195

Total params: 1,610,051

Trainable params: 1,610,051

Non-trainable params: 0

Train on 486 samples, validate on 55 samples

Epoch 1/50

486/486 [=====] - 1s 1ms/step - loss: 2.8259 - acc: 0.7469 - val_loss: 2.8259

Epoch 2/50

486/486 [=====] - 0s 229us/step - loss: 0.5015 - acc: 0.9650 - val_loss: 0.5015

Epoch 3/50

486/486 [=====] - 0s 233us/step - loss: 0.7303 - acc: 0.9403 - val_loss: 0.7303

Epoch 4/50

486/486 [=====] - 0s 223us/step - loss: 0.3563 - acc: 0.9753 - val_loss: 0.3563

Epoch 5/50

486/486 [=====] - 0s 236us/step - loss: 0.3539 - acc: 0.9733 - val_loss: 0.3539

Epoch 6/50

486/486 [=====] - 0s 224us/step - loss: 0.5973 - acc: 0.9609 - val_loss: 0.5973

Epoch 7/50

486/486 [=====] - 0s 238us/step - loss: 0.4245 - acc: 0.9671 - val_loss: 0.4245

Epoch 8/50

486/486 [=====] - 0s 242us/step - loss: 0.3037 - acc: 0.9794 - val_loss: 0.3037

Epoch 9/50

```

486/486 [=====] - 0s 243us/step - loss: 0.2898 - acc: 0.9774 - val_loss: 0.2898
Epoch 10/50
486/486 [=====] - 0s 221us/step - loss: 0.2199 - acc: 0.9856 - val_loss: 0.2199
Epoch 11/50
486/486 [=====] - 0s 218us/step - loss: 0.2319 - acc: 0.9835 - val_loss: 0.2319
Epoch 12/50
486/486 [=====] - 0s 236us/step - loss: 0.4203 - acc: 0.9691 - val_loss: 0.4203
Epoch 13/50
486/486 [=====] - 0s 227us/step - loss: 0.2003 - acc: 0.9877 - val_loss: 0.2003
Epoch 14/50
486/486 [=====] - 0s 224us/step - loss: 0.2901 - acc: 0.9774 - val_loss: 0.2901
Epoch 15/50
486/486 [=====] - 0s 238us/step - loss: 0.3266 - acc: 0.9774 - val_loss: 0.3266
Epoch 16/50
486/486 [=====] - 0s 220us/step - loss: 0.1764 - acc: 0.9877 - val_loss: 0.1764
Epoch 17/50
486/486 [=====] - 0s 222us/step - loss: 0.4236 - acc: 0.9712 - val_loss: 0.4236
Epoch 18/50
486/486 [=====] - 0s 233us/step - loss: 0.1862 - acc: 0.9856 - val_loss: 0.1862
Epoch 19/50
486/486 [=====] - 0s 225us/step - loss: 0.1648 - acc: 0.9897 - val_loss: 0.1648
Epoch 20/50
486/486 [=====] - 0s 216us/step - loss: 0.1048 - acc: 0.9918 - val_loss: 0.1048
Epoch 21/50
486/486 [=====] - 0s 224us/step - loss: 0.2432 - acc: 0.9815 - val_loss: 0.2432
Epoch 22/50
486/486 [=====] - 0s 226us/step - loss: 0.1717 - acc: 0.9856 - val_loss: 0.1717
Epoch 23/50
486/486 [=====] - 0s 232us/step - loss: 0.1682 - acc: 0.9877 - val_loss: 0.1682
Epoch 24/50
486/486 [=====] - 0s 233us/step - loss: 0.2213 - acc: 0.9835 - val_loss: 0.2213
Epoch 25/50
486/486 [=====] - 0s 223us/step - loss: 0.4356 - acc: 0.9712 - val_loss: 0.4356
Epoch 26/50
486/486 [=====] - 0s 221us/step - loss: 0.1415 - acc: 0.9897 - val_loss: 0.1415
Epoch 27/50
486/486 [=====] - 0s 248us/step - loss: 0.1024 - acc: 0.9918 - val_loss: 0.1024
Epoch 28/50
486/486 [=====] - 0s 221us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.1327
Epoch 29/50
486/486 [=====] - 0s 214us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.1327
Epoch 30/50
486/486 [=====] - 0s 228us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.1327
Epoch 31/50
486/486 [=====] - 0s 220us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.1327
Epoch 32/50
486/486 [=====] - 0s 232us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.1327
Epoch 33/50

```

```

486/486 [=====] - 0s 243us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 34/50
486/486 [=====] - 0s 224us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 35/50
486/486 [=====] - 0s 221us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 36/50
486/486 [=====] - 0s 231us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 37/50
486/486 [=====] - 0s 223us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 38/50
486/486 [=====] - 0s 224us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 39/50
486/486 [=====] - 0s 232us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 40/50
486/486 [=====] - 0s 210us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 41/50
486/486 [=====] - 0s 213us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 42/50
486/486 [=====] - 0s 235us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 43/50
486/486 [=====] - 0s 225us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 44/50
486/486 [=====] - 0s 225us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 45/50
486/486 [=====] - 0s 238us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 46/50
486/486 [=====] - 0s 220us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 47/50
486/486 [=====] - 0s 210us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 48/50
486/486 [=====] - 0s 224us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 49/50
486/486 [=====] - 0s 218us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
Epoch 50/50
486/486 [=====] - 0s 223us/step - loss: 0.1327 - acc: 0.9918 - val_loss: 0.2085
541/541 [=====] - 0s 74us/step
['loss', 'acc']
[0.20855218491581207, 0.9870609971599966]

```

```

In [54]: print('VGG 19 - FC - Test accuracy', fc_model.evaluate(features_test, y_test, batch_size=128))
cm = confusion_matrix(np.argmax(fc_model.predict(features_test), axis=1), np.argmax(y_test, axis=1))
df_cm = pd.DataFrame(cm, range(n_classes), range(n_classes))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16})
plt.title('FC model on VGG 19 Features')
plt.show()

```

```

31/31 [=====] - 0s 252us/step

```

VGG 19 - FC - Test accuracy 1.0

